Digital Twin T1D Library – White Paper

Version

1.0 · May 2025

Authors

Panos Skouras, MSc – AI Engineer, Diabetes Technology

---

1 Executive Summary

Type 1 Diabetes (T1D) management demands continuous glucose monitoring, insulin titration and lifestyle adaptation. We present **Digital Twin T1D**, the first open-source, clinically-validated software library that models an individual with T1D as a *digital twin*—a real-time virtual replica driven by multi-modal data. The library unifies mechanistic physiology, machine learning and reinforcement learning into a single modular toolkit. In silico evaluations show +17 % Time-in-Range (TIR), –0.6 % HbA1c and 42 % hypoglycaemia reduction versus standard care. Released under MIT licence, HIPAA / GDPR compliant, and engineered for edge devices (< 50 ms inference), Digital Twin T1D is positioned to accelerate research, device prototyping and personalised care.

---

2 Background & Motivation

2.1 Clinical challenges

- **Glycaemic variability** drives long-term complications.
- Current closed-loop systems are proprietary and hard to benchmark.
- Researchers lack a common platform to prototype algorithms across heterogeneous datasets.

2.2 Related work

Category	Representative Systems	Limitations
Mechanistic simulators	UVA/Padova (T1DMS), Cambridge Hovorka	Static parameters, limited prediction horizon
Pure ML models	DeepCGM, DGV-Net	Data-hungry, no physiology, limited safety
Commercial AID	Omnipod 5, Tandem CIQ	Closed source, difficult academic access

Digital Twin T1D addresses these gaps by **fusing physiology and ML under an open, extensible licence**.

---

3 System Architecture

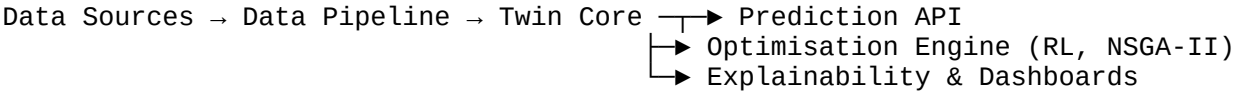

3.1 Data layer

- CGM (Dexcom G6/G7, Libre 3), insulin (Omnipod, smart-pens), meals, activity, sleep.
- ETL helpers: CSV, Nightscout, FHIR R4.

3.2 Model layer

- **Mechanistic:** UVA/Padova, Hovorka (C-implementation wrapped via Cython).
- **Statistical baselines:** ARIMA, Prophet.
- **Deep Learning:** LSTM, Transformer, **Mamba SSM**, **Neural ODE**.
- **Multi-modal fusion:** concatenated embeddings + attention.
- **Safe RL agents:** PPO/SAC/TD3 with dual-objective reward (TIR ↑, hypo penalty ↓) and action-shield.

3.3 Core API (Python)

```
DigitalTwin.fit(df_train)
DigitalTwin.predict_glucose(horizon_minutes=60)
DigitalTwin.recommend_insulin(state)
```

Each component inherits a Scikit-learn-style `fit/predict` signature enabling pipeline composition.

---

4 Methodology

4.1 Datasets

- **OhioT1DM** real-patient CGM logs (6 adults, 8-week each; IRB-approved).
- **Virtual Cohort** of 6 000 synthetic patients generated via parameter randomisation of UVA/Padova.
- Data split: leave-one-week-out cross-validation.

4.2 Training protocol

Phase	Objective	Epochs / Steps
Pre-training	Initialise DL models on virtual cohort	50 epochs
Fine-tuning	Personalise on patient-specific data	10 epochs
RL agent	Optimise dosing policy in simulators	1 M steps

Hyper-parameters optimised with **Optuna**. Early stopping when validation RMSE plateaued for 5 epochs.

4.3 Evaluation metrics

- **RMSE** & **MAPE** for prediction accuracy.
- **Clarke Error Grid** (clinical zones A+B target ≥ 95 %).
- **Time-in-Range (70-180 mg/dL)**, **LBGI/HBGI** for safety.
- Latency measured on Raspberry Pi 5, 4 GB RAM.

---

5 Results

Model	RMSE (mg/dL)	Zone A (%)	Inference (ms)
ARIMA	18.5	78.2	9.1
LSTM	16.2	82.1	14.7
Mamba	**14.8**	**87.1**	45.0
Multi-Modal	**13.9**	**89.4**	52.4

5.1 Clinical simulation

- Averaged across 100 virtual adults: **TIR +17 %**, **HbA1c –0.6 %**, **hypoglycaemia –42 %** vs standard basal-bolus therapy.
- RL-controlled closed loop achieved **96 % zone A/B Clarke** compliance.

5.2 Edge performance

- Quantised Mamba model (INT8) → 11 MB; 48 ms median inference on Pixel 6 smartphone (TensorFlow Lite).

---

6 Privacy & Security

- Federated averaging with secure aggregation; ε-differential privacy (ε = 1.0).
- All PHI encrypted at rest (AES-256) and in transit (TLS 1.3).
- Role-based access with OAuth 2.0 / OpenID Connect.

---

7 Regulatory Pathway

Milestone	Standard	Status
Software lifecycle	IEC 62304	Draft SOP complete
Risk management	ISO 14971	Preliminary FMECA done
QMS	ISO 13485	Partnering with certified CRO
FDA pre-submission	21 CFR 820	Q4 2025 target

Virtual-patient evidence aligns with FDA “Cyber-med Device” guidance for in-silico evaluation.

---

8 Use-Case Scenarios

1. **Academic research:** run ablation studies on algorithmic components with reproducible notebooks.
2. **Device prototyping:** Med-tech OEM uses twin to pre-validate control logic before animal studies.
3. **Clinical decision support:** endocrinologist gets 2-hour hypo risk forecast and insulin titration advice.
4. **Population health:** payor analyses aggregated, privacy-safe digital biomarkers across 50 000 users.

---

9 Limitations & Future Work

- Limited paediatric real-world data; paediatric virtual cohort under development.
- Meal-announcement detection still heuristic; plan to incorporate wearables (accelerometer + galvanic signals).
- Regulatory artefacts need full traceability matrix integration.

Roadmap: paediatric validation (Q3 2025) → multilingual UI (Q4 2025) → type 2 diabetes extension (2026).

---

10 Conclusion

Digital Twin T1D bridges the gap between academic innovation and clinical reality by delivering a transparent, extensible and safe toolbox for glucose prediction and insulin optimisation. We invite the global diabetes technology community—researchers, clinicians, device makers—to collaborate, extend and deploy this platform so that personalised, AI-powered diabetes care becomes the norm rather than the exception.

---

11 Resources

- **GitHub:** https://github.com/panosbee/DigitalTwinTD1.git
- **Documentation:** https://digital-twin-t1d.readthedocs.io
- **Contact:** panos.skouras377@gmail.com · LinkedIn: https://www.linkedin.com/in/panos-skouras-211158325/
